

Implementing a proof-of-concept cloud architecture

CLOUD SERVICES

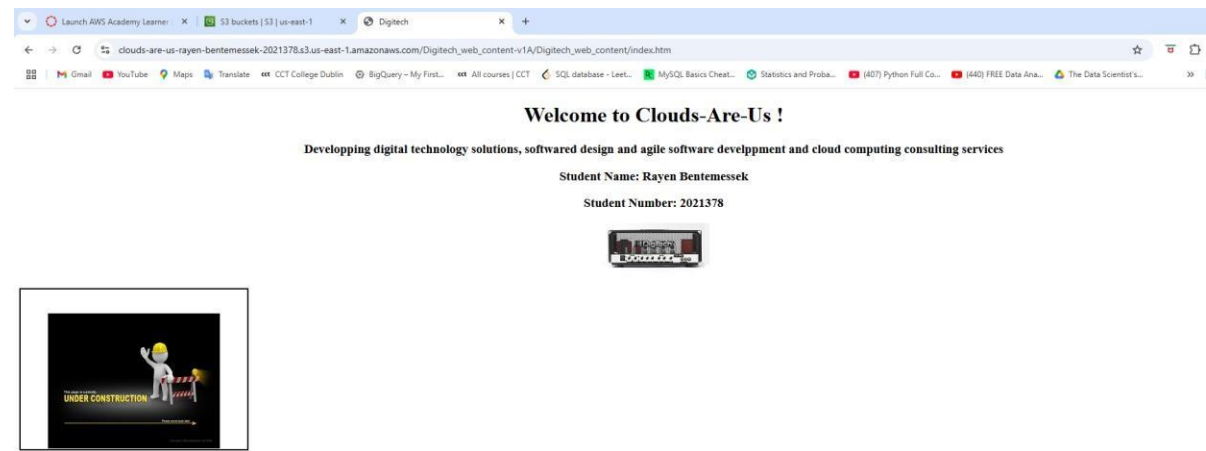
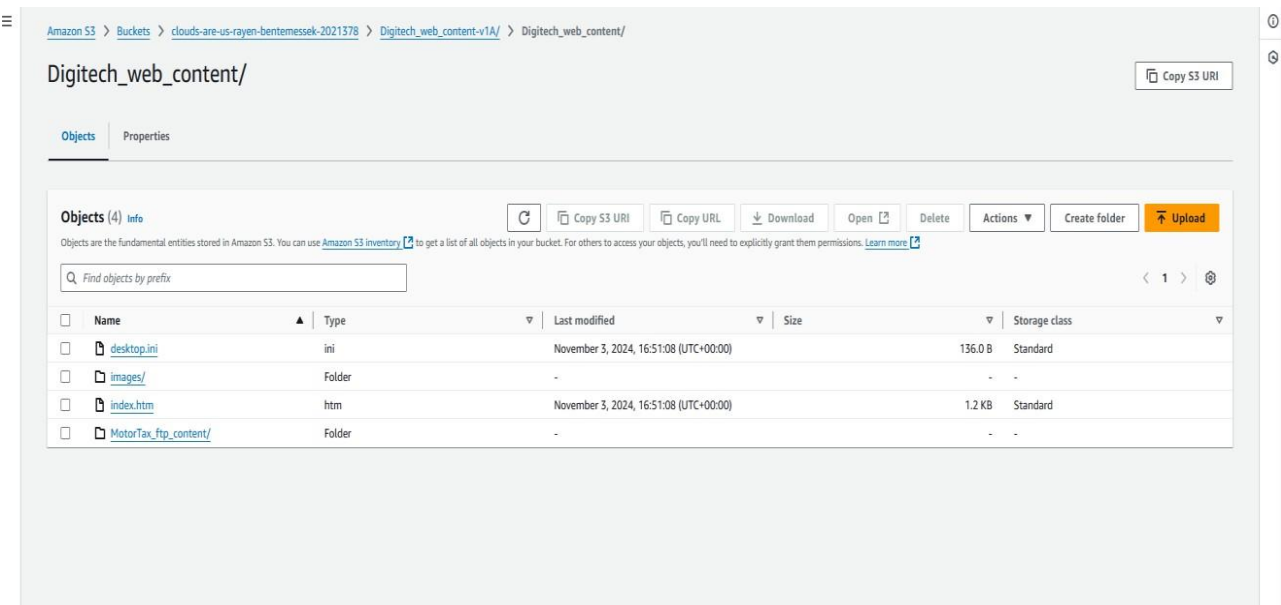
RAYEN BENTEMESSEK

Table of Contents

Task 1:.....	3
Task 1a:	3
Task 1b:.....	4
Task 2:.....	6
Task 2a:	6
Task 2b:.....	8
Task 3:.....	10
Task 3a:	10
Task 3b:.....	13
References	14

Task 1:

Task 1a:



Task 1b:

AWS S3 is a simple storage service by Amazon useful for hosting website resources such as images and videos and more. S3 is an object-level data storage which means it manages data as objects and stores them in a flat data environment where they are distributed across different machines and allows users to access them using internet from anywhere in the world (GeeksforGeeks, 2021).

AWS EBS, elastic block store, is a block-level storage that stores data files on storage area networks (SAN) in multiple volumes called blocks, these blocks are separate hard drives, and the storage cannot be accessible through the internet unlike the S3 (GeeksforGeeks, 2021).

Key differences:

Characteristic	AWS S3	AWS EBS
Storage type	Object storage	Block storage
Accessibility	Can be retrieved using HTTP or BitTorrent (internet accessible)	Only accessed through the instance connected to
Access control	Access base on IAM, uses bucket policy and can control public access and add restrictions if needed	Security groups and Identity and access management (IAM)
Availability	Accessed through API's	Accessed by a single instance
Durability	Multiple AZ redundancy and highly durable	Single AZ and highly durable
Storage and data size limits	No limits on the number of objects, each object size can be up to 5TB	Maximum storage size of 16 TB, no file size limit on disk
Pricing	Free Tier: 5 GB First 50 TB/month: \$0.023/GB 450 TB/month: \$0.022/GB Over 500 TB / Month: \$0.021 per GB	Free Tier: 30 GB, 1 GB snapshot storage General Purpose: \$0.045/GB(1 month) Provisioned SSD: \$0.125/GB(1 month)
Scalability	Can be automatically scalable where resources can be provisioned and de-provisioned	Manual scalability to increase or decrease storage resources

Source: (GeeksforGeeks, 2021); (Fayaz, 2023)

Similarities:

- **Security:** Both S3 and EBS support data encryption at rest and in transmission which makes them both highly secure (GeeksforGeeks, 2021).
- **Availability:** Despite having different implementations of availability zones (multi-AZ vs Single AZ) but they both have offer (GeeksforGeeks, 2021).

Use cases:

Amazon S3:


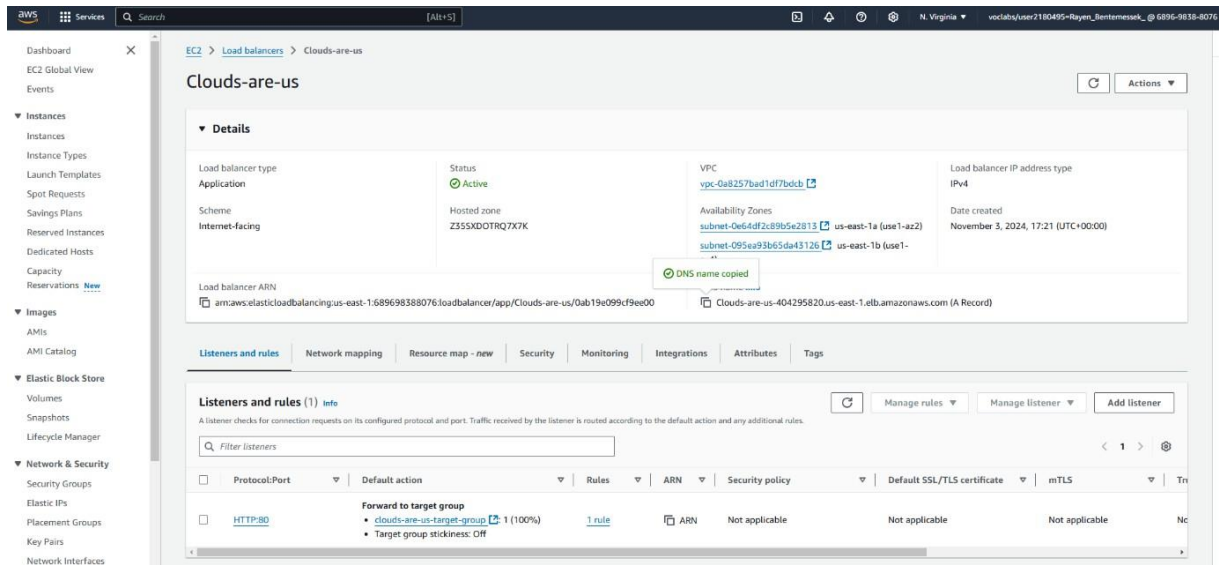
- **Data lake and big data analytics:** Amazon S3 works with AWS Lake formation which will create data lakes that hold raw data in native format and then support big data analytics using machine learning tools, query-in-place and finally draws important insights from the raw data (GeeksforGeeks, 2021).
- **Backup and restoration:** Amazon S3 combined with other services (EBS, EFS, etc) can provide a reliable and robust backup solution (GeeksforGeeks, 2021).
- **Disaster recovery:** S3 provides reliable and powerful data recovery options from different types of disasters such as power outages, systems failures and human errors (GeeksforGeeks, 2021).

Amazon EBS:

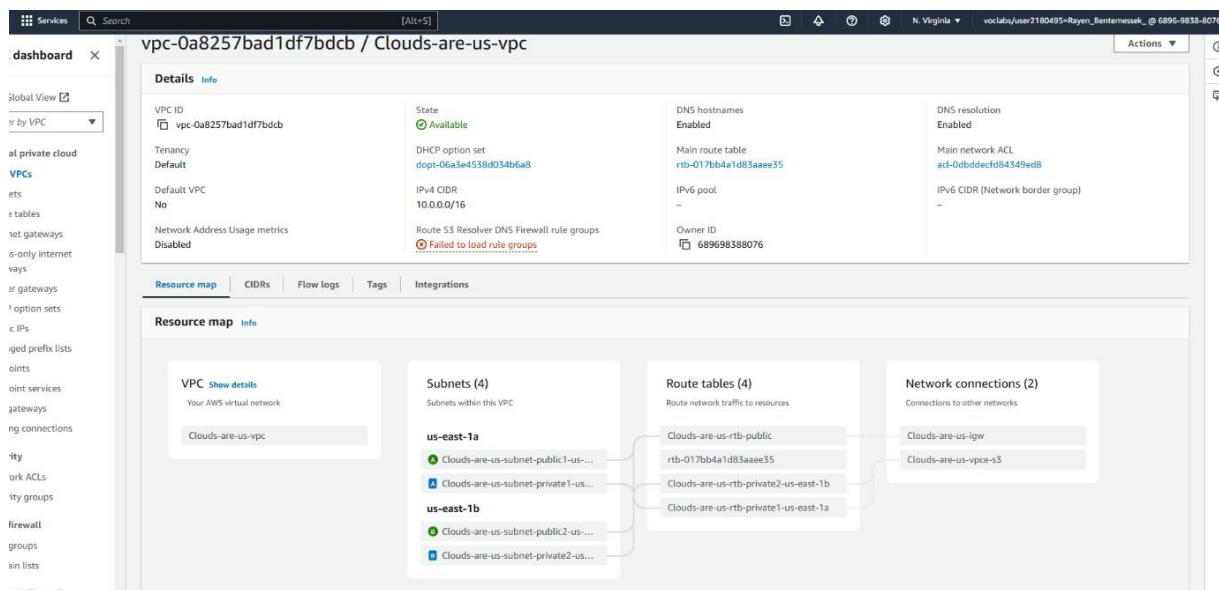
- **Software Testing and development:** Amazon EBS is ideal for applications as it connected only to a particular instance which helps in reducing latency and maintains storage in one location (GeeksforGeeks, 2021).
- **Business Continuity:** Amazon EBS provides a high level of consistency throughout different AWS region which can make data persistent making it an ideal solution for businesses to rely on It especially when they are looking for data persistence (GeeksforGeeks, 2021).
- **Transactional and NoSQL databases:** EBS provides low level of latency which offers a great performance making it a great option for high-performance transactional databases and in database management (GeeksforGeeks, 2021)

Task 2:

Task 2a:



DigiTech! This is Server ip-10-0-128-147.ec2.internal



This picture shows the implementation of a custom VP

Launch AWS Academy Learner | Target group details | EC2 | us-east-1 | Instances | EC2 | us-east-1 | Digitech | Launch an instance | EC2 | us-east-1 | Digitech_web_content/images | Digitech_web_content/images | us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#instances:instanceState=running

Services Search [Alt+S]

Dashboard EC2 Global View Events

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Capacity

Reservations **New**

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Instances (1/5) Info

Find instance by attribute or tag (case-sensitive)

Instance state = running Clear filters

All states

Last updated 3 minutes ago Connect Instance state Actions Launch instances

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IP
<input checked="" type="checkbox"/> Clouds-are-us	i-086bd3226bad27f28	Running	t2.micro	Initializing	View alarms +	us-east-1a	ec2-3-85-135-21.comp...	3.85.135.
<input type="checkbox"/> Clouds-are-us	i-0c9c99a518f9b5e75	Running	t2.micro	Initializing	View alarms +	us-east-1a	ec2-18-206-231-64.co...	18.206.2.
<input type="checkbox"/> Clouds-are-us	i-042e6a12bcb1cb0b0	Running	t2.micro	Initializing	View alarms +	us-east-1a	ec2-3-83-40-154.comp...	3.83.40.1
<input type="checkbox"/> Clouds-are-us	i-0ff6c39e6d433acf	Running	t2.micro	Initializing	View alarms +	us-east-1a	ec2-34-226-152-7.com...	34.226.1
<input type="checkbox"/> Clouds-are-us	i-0d0c9fe3e3d6b1f35	Running	t2.micro	Initializing	View alarms +	us-east-1a	ec2-52-91-234-139.co...	52.91.23.

i-086bd3226bad27f28 (Clouds-are-us)

Details Status and alarms Monitoring Security Networking Storage Tags

Instance summary Info

Instance ID

i-086bd3226bad27f28

Public IPv4 address

3.85.135.21 | open address

IPv6 address

-

Instance state

Running

Private IP address

ip-10-0-2-194.ec2.internal

Private IP DNS name (IPv4 only)

ip-10-0-2-194.ec2.internal

Instance type

t2.micro

VPC ID

vpc-0a8257bad1df7bdc (Clouds-are-us-vpc)

Subnet ID

subnet-0e64d7c89b5e2813 (Clouds-are-us-subnet-public1-us-east-1a)

IAM Role

LabRole

Auto-assigned IP address

3.85.135.21 [Public IP]

Private IPv4 addresses

10.0.2.194

Public IPv4 DNS

ec2-3-85-135-21.compute-1.amazonaws.com | open address

Elastic IP addresses

-

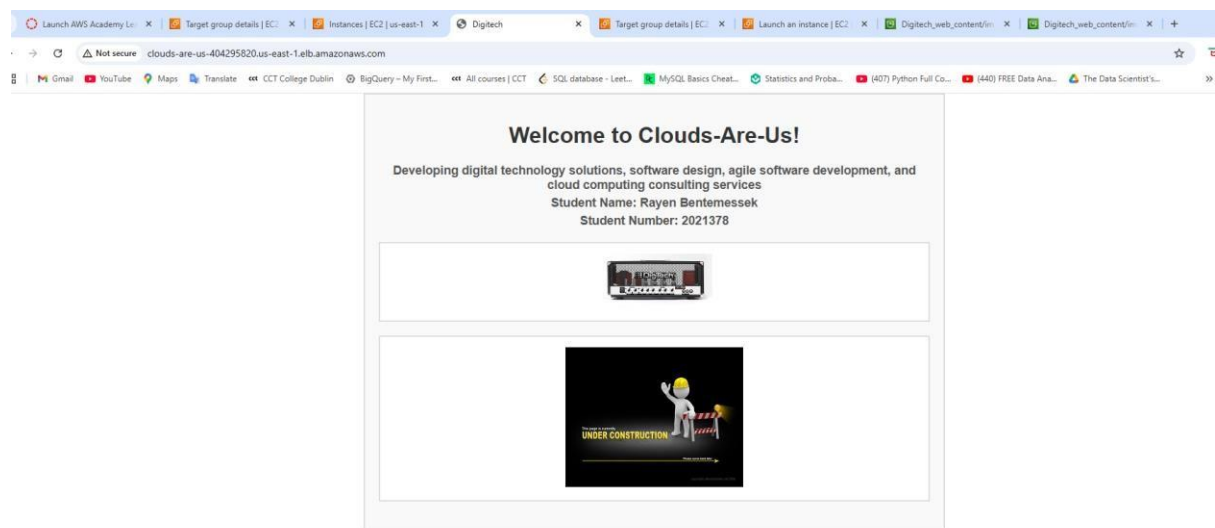
AWS Compute Optimizer finding

Opt-in to AWS Compute Optimizer for recommendations. | Learn more

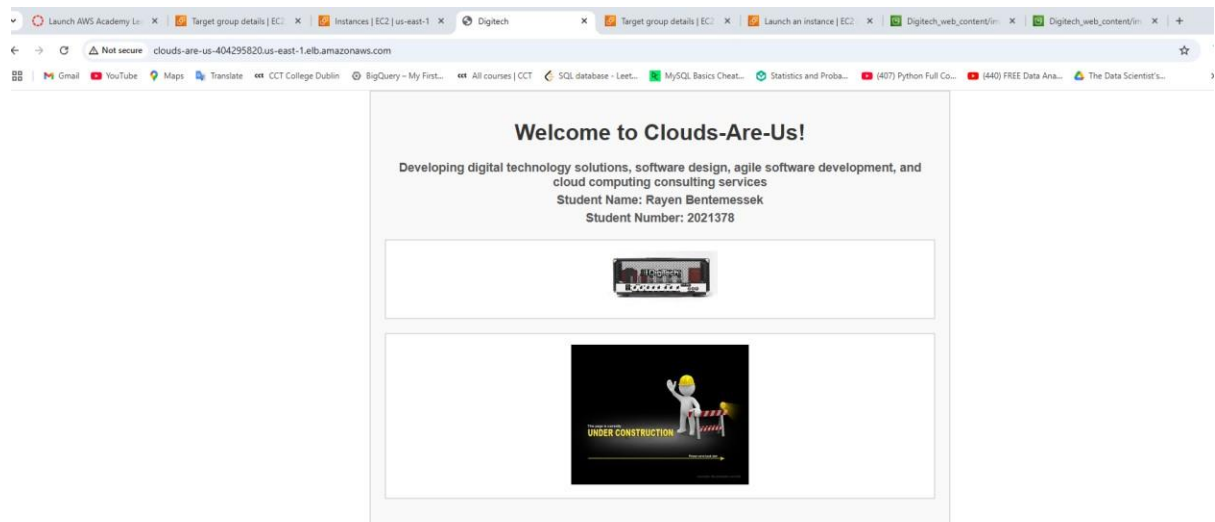
Auto Scaling Group name

-

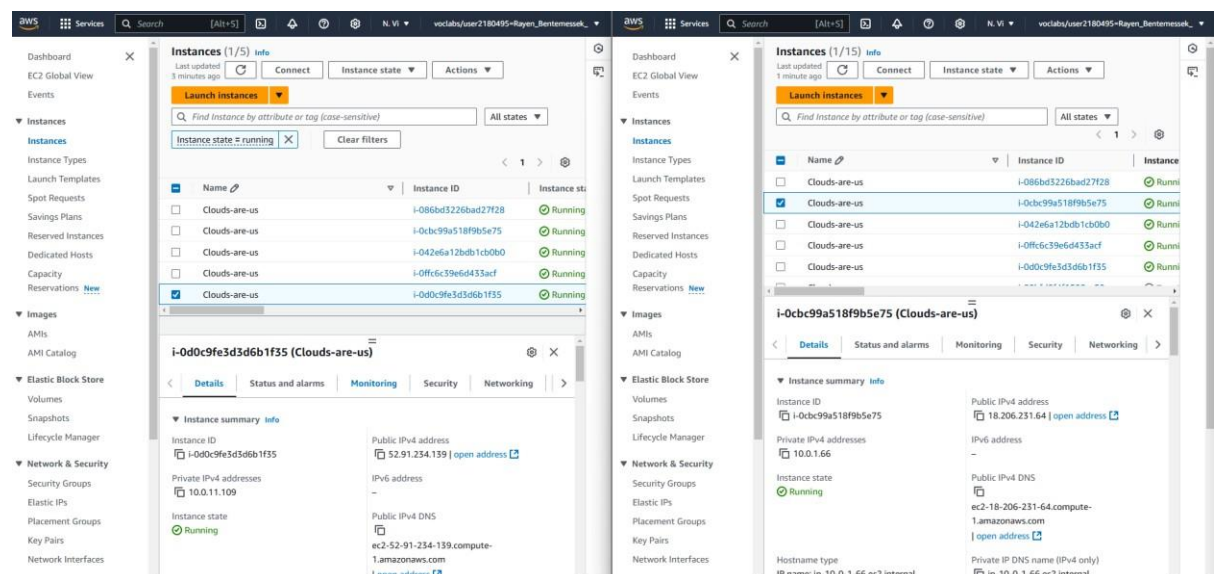
Task 2b:

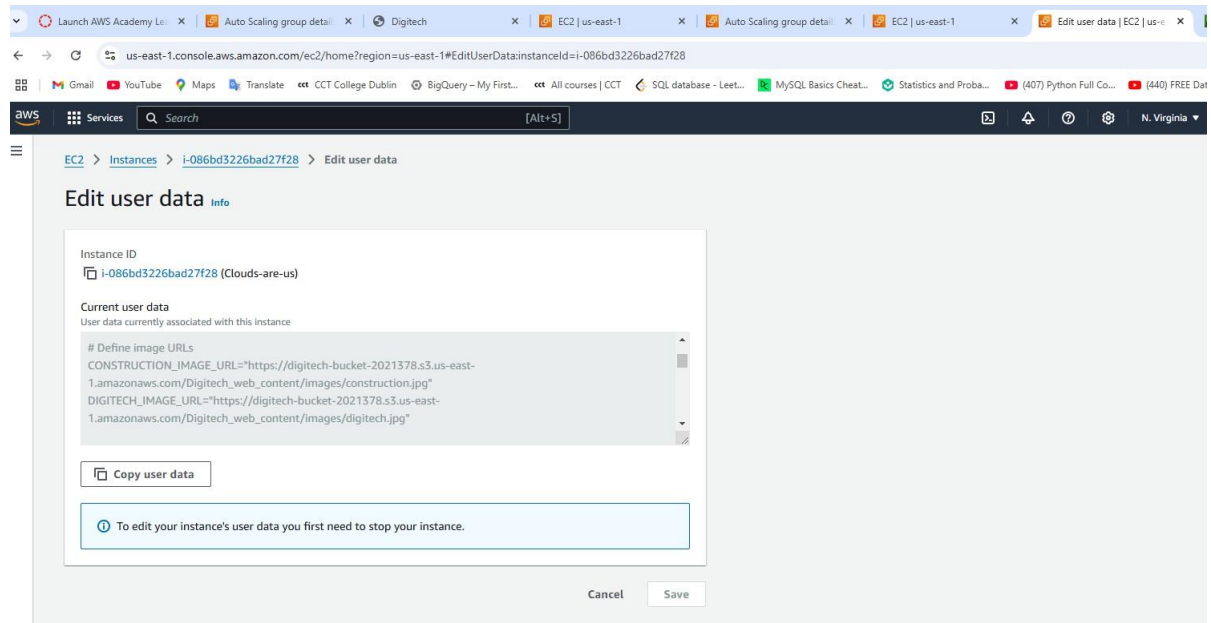


This is Server ip-10-0-11-109.ec2.internal



This is Server ip-10-0-1-66.ec2.internal





This task was a little bit tricky because I had too many issues in the beginning as I was trying to access the whole website from the S3 bucket using a script but that didn't work because of IAM roles permissions, so what I did is create the HTML page in the script and only access the images through their S3 URLs as shown in the screenshot above. This worked eventually and I managed to display the desired websites across the five instances showing each unique IP address using a custom VPC called clouds-are-us as shown in the task 2a.

Task 3:

Task 3a:

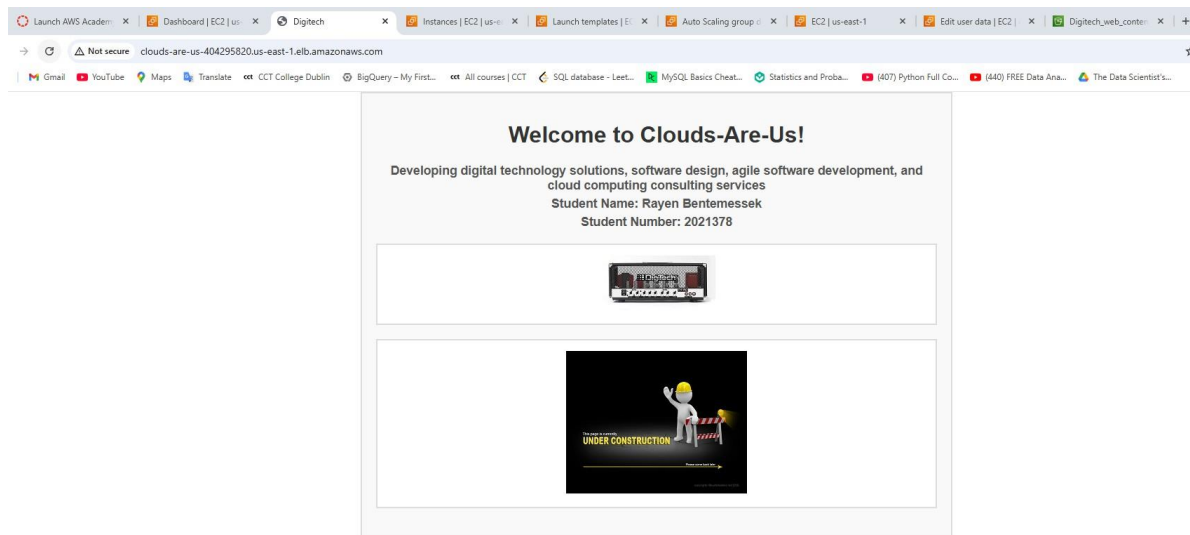
The screenshot displays the AWS Management Console interface. The top navigation bar shows the 'Auto Scaling group details' page for the 'digitech-asg' group in the 'us-east-1' region. The left sidebar contains the 'Instances' section, which is currently expanded to show 'Auto Scaling Groups'. The main content area is divided into several sections: 'Group details', 'Launch template', 'Network', 'Instance type requirements', and 'Load balancing'. The 'Group details' section shows the group name 'digitech-asg', the desired capacity of 2, and the Amazon Resource Name (ARN). The 'Launch template' section shows the template name 'digitech', the AMI ID 'ami-063d43db0594b521b', and the instance type 't2.micro'. The 'Network' section shows the availability zones 'us-east-1b' and 'us-east-1c', and the subnet ID 'subnet-095ea93b65da43126'. The 'Instance type requirements' section shows the instance type 't2.micro'. The 'Load balancing' section shows the load balancer target group 'clouds-are-us-target-group'. Below these sections, a table lists the instances in the group. The table has columns for Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, Public IPv4 DNS, and Public IP. Two instances are listed: 'i-017f64814a3121a6e' and 'i-04550397bfc16b22f', both in a 'Running' state. Below the table, the details for the instance 'i-017f64814a3121a6e' are shown, including its Instance ID, Public IPv4 address, Instance state, Private IP DNS name, Instance type, VPC ID, Subnet ID, Instance ARN, Platform, and Monitoring status.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IP
i-017f64814a3121a6e	i-017f64814a3121a6e	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-3-91-66-168.comp...	3.91.66.1
i-04550397bfc16b22f	i-04550397bfc16b22f	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	ec2-18-206-174-15.co...	18.206.174.15

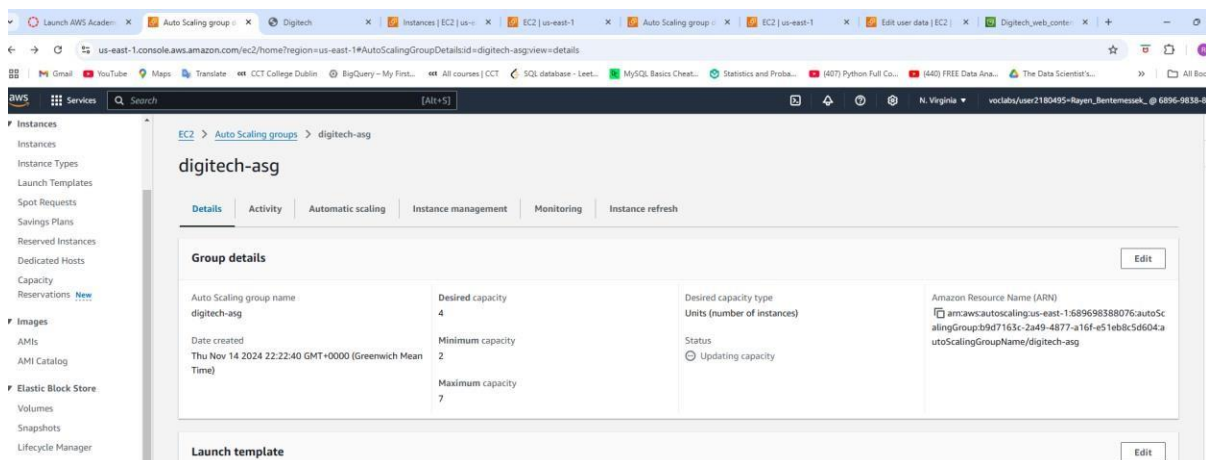
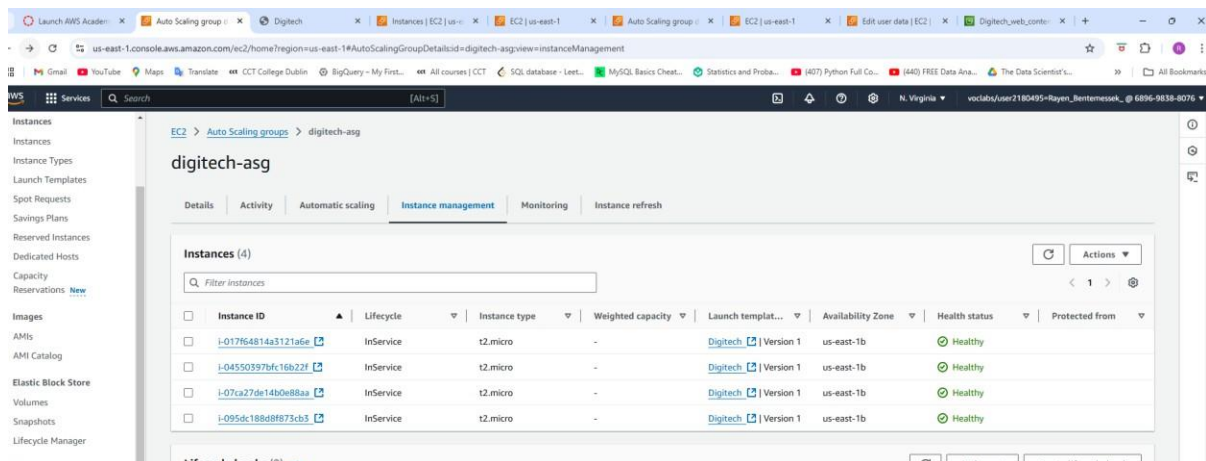
Instance details for i-017f64814a3121a6e:

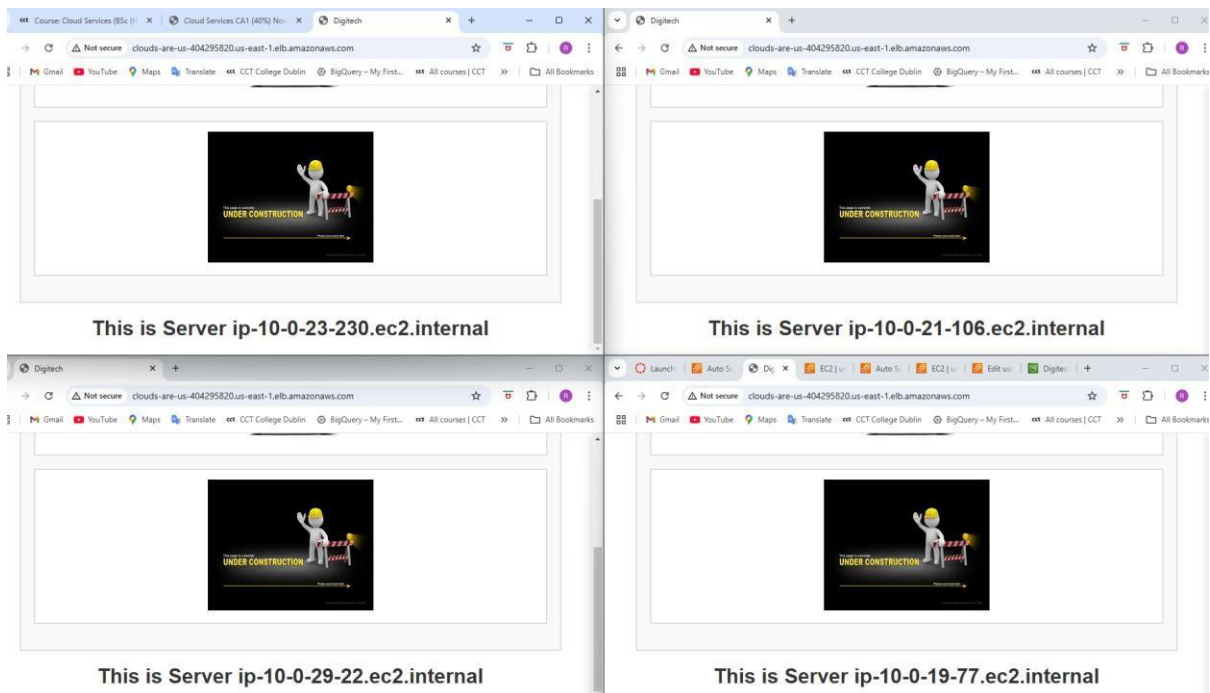
- Instance ID: i-017f64814a3121a6e
- Public IPv4 address: 3.91.66.168 [open address]
- Instance state: Running
- Private IP DNS name (IPv4 only): ip-10-0-23-230.ec2.internal
- Instance type: t2.micro
- VPC ID: vpc-0a8257bad1df7dbcb [Clouds-are-us-vpc]
- Subnet ID: subnet-095ea93b65da43126 [Clouds-are-us-subnet-public2-us-east-1b]
- Instance ARN: arn:aws:ec2:us-east-1:689698388076:instance/i-017f64814a3121a6e
- Platform: Amazon Linux
- Monitoring: disabled

I created the ASG, used the same script in task 2b in the launch template, I created it with a desired capacity of 2 instances at first and later I have manually scaled up and changed desired capacity to 4 instances to show that the ASG is working perfectly for what we need it for.



This is Server ip-10-0-23-230.ec2.internal





Benefits of integrating ASG combined with the ALB:

In this scenario, the ASG will automatically adjust the number of instances based on the current demand based on the initial configuration; the more traffic comes, the more machines the ASG will create and vice versa. The ALB will adapt to the current number of instances and distribute the traffic across the available healthy instances, ensuring each handles a reasonable amount of traffic without overusing it. If one machine does not pass the health check, the ALB will not use this instance until it is automatically (or manually) replaced by the ASG. This approach is very common in the real world; for example, an online retail company would use this combination as part of their architecture, ensuring that their application is highly scalable during peak times (Christmas, for example) to handle the unpredictable amount of traffic coming, ensuring high availability during this period. As a result, the company will make more profits by simply reducing the costs of servers running to replace them with dynamically allocated servers and optimizing their website presence, which will create more sales. Also, this will help update the website as they will introduce new offers and deals; the website will not be down when these updates are implemented. This makes integrating ASG with an ALB highly beneficial for those who use it.

Task 3b:

The DigiTech Application Load Balancer has plenty of benefits for the DigiTech application as it distributes incoming traffic across multiple application servers as long as they are available, making it essential for the application because this ensures that no single server handles too much traffic. In fact, this is used even in real life by providing different routes to reach a certain place, which will distribute the traffic (cars) among all the roads making everybody reach on time, if we only have one major route, all the cars will choose it and there will be a huge traffic jam. What makes an ALB even better is that the routes can be increased or decreased based on demands if it's combined with an Auto scaling group, so if there's a lot of demand like in

Christmas time for shopping websites, there will be as much routes as possible (based on the initial configuration) and when there's not much traffic these routes will be reduced to save costs and unused resources.

This will ensure many benefits for DigiTech such as:

- **Availability and Scalability:** As explained earlier on, this setup will ensure that the DigiTech application is highly available based on the user demands by providing enough resources to satisfy the demand, which is done by the ASG, while the ALB will ensure to divide the incoming traffic based on the resources provided by the ASG. As a result, DigiTech will always be available and provide a smooth user experience for everyone ensuring the business continuity.
- **Fault Tolerance:** The ALB, apart from the previous task, will make sure users are redirected to a healthy instance by constantly passing the instances through health check, which will let users keep using the website, and give developers enough time to fix any issues related to faulty or unhealthy instances.

By combining the ALB with the ASG, DigiTech will manage to follow some of the AWS Well- Architected framework, such as:

- **Reliability Pillar:** This Pillar focuses on creating a reliable framework that can automatically recover from failure by automatically notifying and tracking failures, also it focuses on scaling horizontally by creating multiple small resources to reduce the impact of a single failure (We talked about this in the Previous homework) and distribute requests across multiple resources and also

stop guessing capacity by automatically allocating resources based on demands. As explained above, the combination of an ASG with an ALB will make sure this pillar is achieved which is highly crucial (docs.aws.amazon.com, n.d.).

- **Cost optimization Pillar:** Every business looks for reducing costs, and spending money on the cloud can be highly expensive if not planned carefully, that's why AWS highlights this pillar, which consists of reducing costs while maintaining high performance. The ASG will make sure that only needed resources are existing, The only criteria to follow here is the traffic, so during low-traffic periods the ASG will try to reduce the resources as much as possible while the ALB will distribute the traffic among all the resources, however in peak times, where most traffic usually happens we would expect the ASG to maximize the resources until the traffic goes down, this approach will reduce the cost of running instances following the model pay-as-go where the business only pays for the used resources which can help save tons of money, and this remains one of the most important reasons companies are shifting from on-premise to the cloud resources as it makes a difference over time (docs.aws.amazon.com, n.d.).

References:

GeeksforGeeks. (2021). Difference Between AWS S3 and AWS EBS. [online] Available at: <https://www.geeksforgeeks.org/difference-between-aws-s3-and-aws-ebs/>.

Fayaz, M. (2023). Overview of AWS EFS, EBS and S3 and their differences. [online] Cloud Training Program. Available at: https://k21academy.com/amazon-web-services/difference-between-aws-efs-ebs-and-s3/#Improves_query_performance_by_using_S3_Select_to_retrieve_subsets_of_object_data_in stead_of_the_entire_object.

docs.aws.amazon.com. (n.d.). Design Principles - Reliability Pillar. [online] Available at: <https://docs.aws.amazon.com/wellarchitected/latest/reliability-pillar/design-principles.html>.

docs.aws.amazon.com. (n.d.). Design Principles - AWS Well-Architected Framework. [online] Available at: <https://docs.aws.amazon.com/wellarchitected/latest/framework/cost-dp.html>.